



Jade Yun, Si Chen, Yimei Chen, Yu Tian

Data

kaggle

Search kaggle

Q

Competitions

Reviewed Dataset

SF Bay Area Bike Share

Anonymized bike trip data from August 2013 to August 2015

Ben Hamner

last updated 2 years ago

Overview

Data

Kernels

Discussion

Activity

Tags

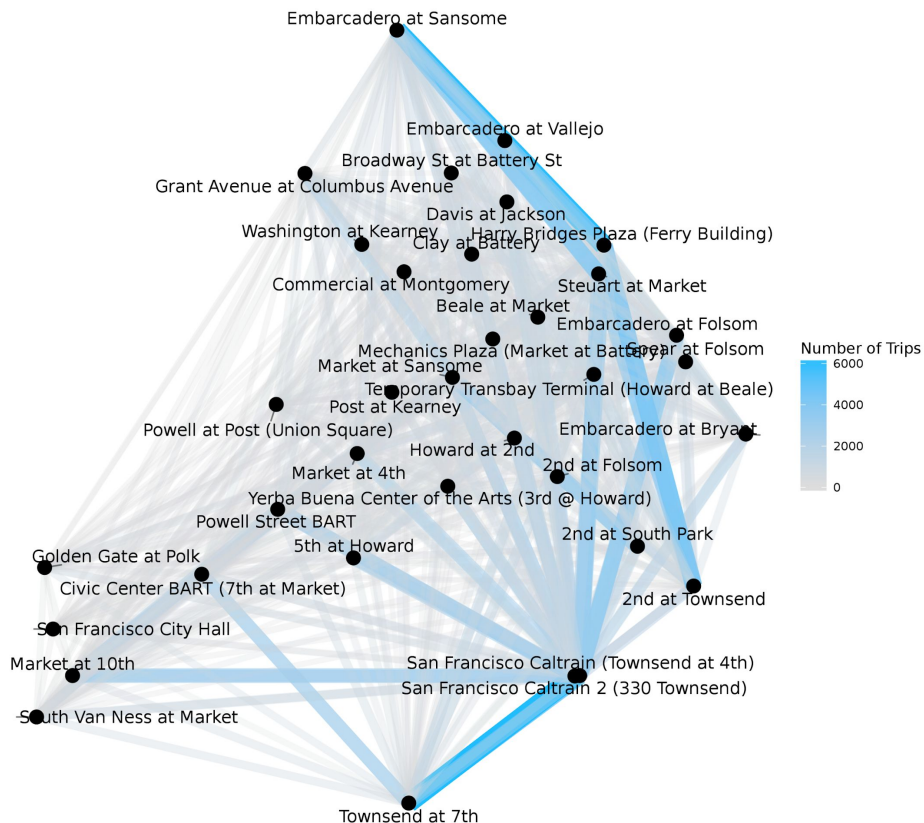
road transport

cycling

large

featured

Data Size: 2.71GB
Format: CSV
Time Window: 2013/8 - 2015/8 (2 yrs)



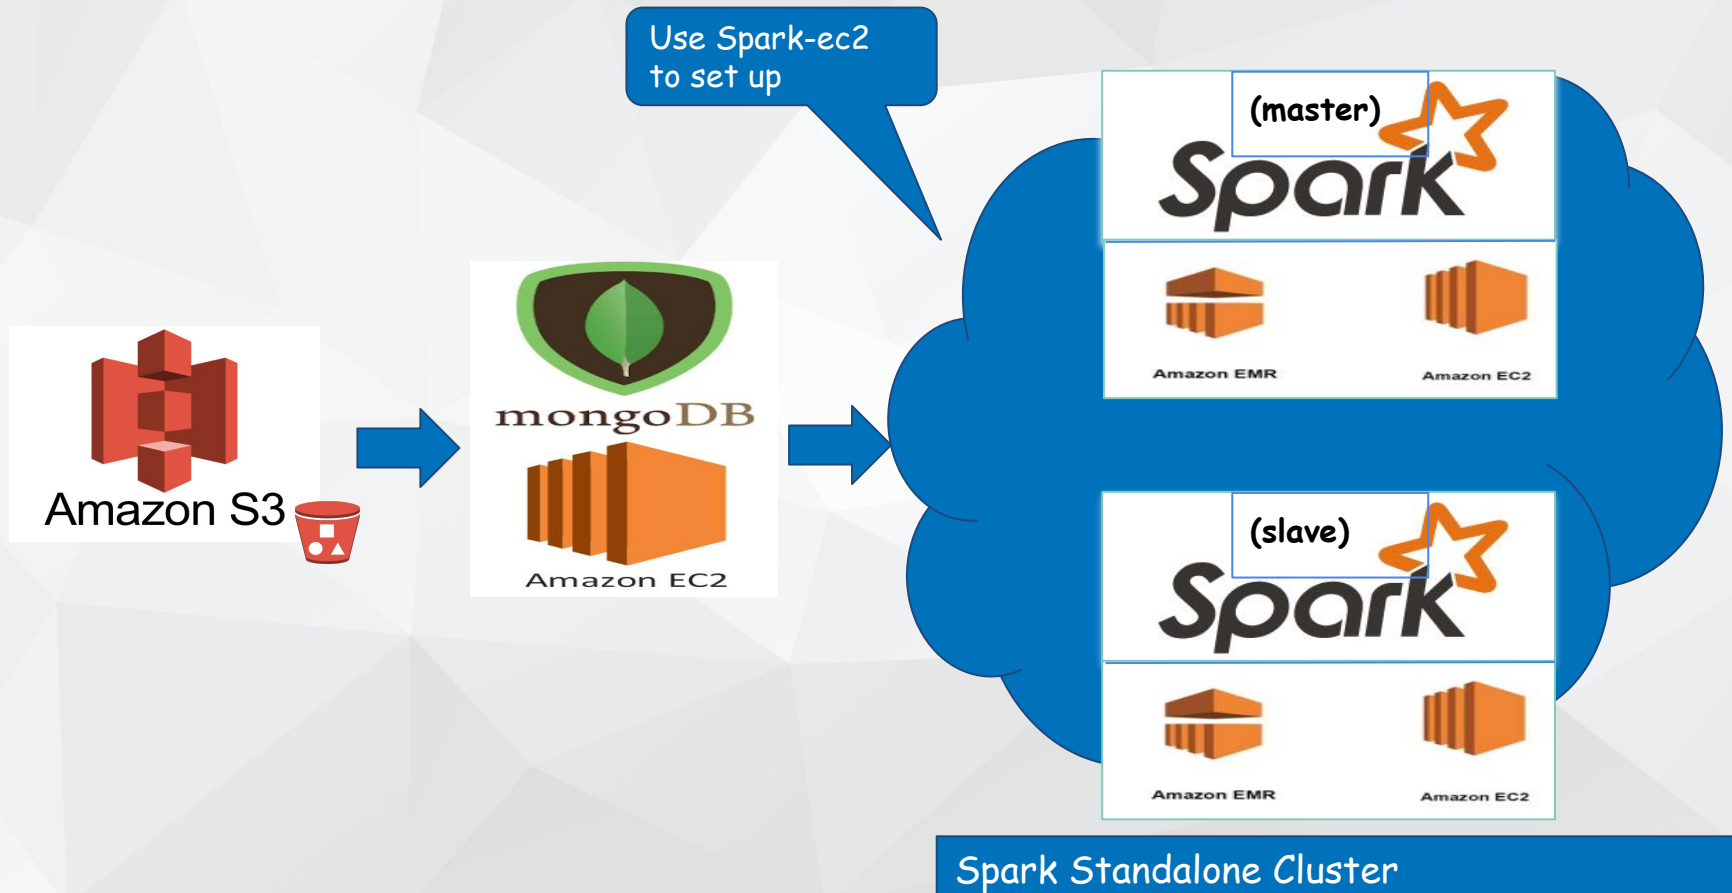
Why choose this data?

- Provide convenience to big city commute
- A lot of insights can be generated
- Interesting topic: smart city

Goal?

Predicting daily demand for each station

The Pipeline We Use



1. S3 - Upload Data

Amazon S3 > sparkgroup123

Overview Properties Permissions Public Management

Q Type a prefix and press Enter to search. Press ESC to clear.

Upload Create folder More

US West (Oregon) ↻

Viewing 1 to 4

<input type="checkbox"/>	Name ↑	Last modified ↑	Size ↑	Storage class ↑
<input type="checkbox"/>	station.csv	Jan 13, 2018 2:55:46 PM GMT-0800	5.5 KB	Standard
<input type="checkbox"/>	status.csv	Jan 13, 2018 2:56:11 PM GMT-0800	1.9 GB	Standard
<input type="checkbox"/>	trip.csv	Jan 13, 2018 3:05:19 PM GMT-0800	76.5 MB	Standard
<input type="checkbox"/>	weather.csv	Jan 13, 2018 2:55:46 PM GMT-0800	427.8 KB	Standard

For larger data

Build an EC2 to download data and use command line to upload to S3

2. MongoDB Server (t2.medium)

1. Download data from S3

```
[ec2-user@ip-172-31-23-234 ~]$ aws s3 cp s3://sparkgroup123/status.csv .
download: s3://sparkgroup123/status.csv to ./status.csv
[ec2-user@ip-172-31-23-234 ~]$ aws s3 cp s3://sparkgroup123/trip.csv .
download: s3://sparkgroup123/trip.csv to ./trip.csv
[ec2-user@ip-172-31-23-234 ~]$ aws s3 cp s3://sparkgroup123/weather.csv .
download: s3://sparkgroup123/weather.csv to ./weather.csv
[ec2-user@ip-172-31-23-234 ~]$
```

Model	vCPU	CPU Credits / hour	Mem (GiB)
t2.medium	2	24	4

2. Load data into MongoDB

```
[ec2-user@ip-172-31-23-234 ~]$ mongoimport --db bikeshare --collection status --file status.csv
--type csv --headerline
2018-01-15T02:58:06.481+0000    connected to: localhost
2018-01-15T02:58:09.480+0000    [.....] bikeshare.status    9.26MB/1.85GB (0.5%)
2018-01-15T02:58:12.480+0000    [.....] bikeshare.status    18.1MB/1.85GB (1.0%)
2018-01-15T02:58:15.483+0000    [.....] bikeshare.status    27.2MB/1.85GB (1.4%)
2018-01-15T02:58:18.480+0000    [.....] bikeshare.status    36.2MB/1.85GB (1.9%)
2018-01-15T02:58:21.480+0000    [.....] bikeshare.status    45.0MB/1.85GB (2.4%)
2018-01-15T02:58:24.480+0000    [.....] bikeshare.status    54.0MB/1.85GB (2.8%)
2018-01-15T02:58:27.481+0000    [.....] bikeshare.status    63.3MB/1.85GB (3.3%)
```

3. Processing & Query

```
> db.station.findOne()
{
  "_id" : ObjectId("5a5d461db2c867959cfd1198"),
  "id" : 3,
  "name" : "San Jose Civic Center",
  "lat" : 37.330698,
  "long" : -121.888979,
  "dock_count" : 15,
  "city" : "San Jose",
  "installation_date" : "8/5/2013"
}
> db.station.count()
70
> db.station.update({},{$rename:{'id':'station_id'}},{multi:true})
WriteResult({"nMatched" : 70, "nUpserted" : 0, "nModified" : 70 })
```


What we learnt from the pipeline

Connect from my laptop to that mongodb to check data

t2.micro/t2.small: the importing status.csv (the big data) quits in the middle!!!

```
[ec2-user@ip-172-31-27-156 ~]$ mongoimport --db bikeshare --collection status --file status.csv --type csv --headerline
connected to: localhost
2018-01-17T20:51:12.861+0000 [.....] bikeshare.status 6.27MB/1.85GB (0.3%)
2018-01-17T20:51:15.861+0000 [.....] bikeshare.status 12.2MB/1.85GB (0.6%)
2018-01-17T20:51:18.873+0000 [.....] bikeshare.status 18.0MB/1.85GB (0.9%)
2018-01-17T20:51:21.874+0000 [.....] bikeshare.status 23.7MB/1.85GB (1.2%)
2018-01-17T20:51:24.864+0000 [.....] bikeshare.status 29.6MB/1.85GB (1.6%)
2018-01-17T20:51:27.861+0000 [.....] bikeshare.status 35.4MB/1.85GB (1.9%)
2018-01-17T20:51:30.863+0000 [.....] bikeshare.status 41.2MB/1.85GB (2.2%)
2018-01-17T20:51:33.868+0000 [.....] bikeshare.status 46.9MB/1.85GB (2.5%)
2018-01-17T20:51:36.861+0000 [.....] bikeshare.status 52.5MB/1.85GB (2.8%)
2018-01-17T20:51:39.861+0000 [.....] bikeshare.status 58.6MB/1.85GB (3.1%)
2018-01-17T20:51:42.864+0000 [.....] bikeshare.status 64.7MB/1.85GB (3.4%)
2018-01-17T20:51:45.861+0000 [.....] bikeshare.status 70.6MB/1.85GB (3.7%)
2018-01-17T20:51:48.865+0000 [.....] bikeshare.status 76.5MB/1.85GB (4.0%)
2018-01-17T20:51:51.861+0000 [.....] bikeshare.status 82.6MB/1.85GB (4.4%)
2018-01-17T20:51:54.867+0000 [#.....] bikeshare.status 88.5MB/1.85GB (4.7%)
2018-01-17T20:51:57.861+0000 [#.....] bikeshare.status 94.5MB/1.85GB (5.0%)
2018-01-17T20:52:00.864+0000 [#.....] bikeshare.status 100MB/1.85GB (5.3%)
2018-01-17T20:52:03.868+0000 [#.....] bikeshare.status 107MB/1.85GB (5.6%)
2018-01-17T20:52:06.870+0000 [#.....] bikeshare.status 112MB/1.85GB (5.9%)
2018-01-17T20:52:09.861+0000 [#.....] bikeshare.status 118MB/1.85GB (6.2%)
2018-01-17T20:52:12.869+0000 [#.....] bikeshare.status 124MB/1.85GB (6.5%)
2018-01-17T20:52:15.861+0000 [#.....] bikeshare.status 130MB/1.85GB (6.8%)
2018-01-17T20:52:18.861+0000 [#.....] bikeshare.status 135MB/1.85GB (7.1%)
2018-01-17T20:52:21.862+0000 [#.....] bikeshare.status 137MB/1.85GB (7.2%)
2018-01-17T20:52:22.759+0000 [#.....] bikeshare.status
2018-01-17T20:52:22.760+0000 Failed: lost connection to server
2018-01-17T20:52:22.760+0000 imported 5281000 documents
```

3. Spark - Load Data from MongoDB

```
weather= spark.read.format('com.mongodb.spark.sql.DefaultSource').option('uri','mongodb://54.202.150.222/bikeshare.weather')
trip= spark.read.format('com.mongodb.spark.sql.DefaultSource').option('uri','mongodb://54.202.150.222/bikeshare.trip')
station= spark.read.format('com.mongodb.spark.sql.DefaultSource').option('uri','mongodb://54.202.150.222/bikeshare.station')
status= spark.read.format('com.mongodb.spark.sql.DefaultSource').option('uri','mongodb://54.202.150.222/bikeshare.status')
```

```
station=station.drop('_id')
station.printSchema()
```

```
root
|-- city: string (nullable = true)
|-- dock_count: integer (nullable = true)
|-- installation_date: string (nullable = true)
|-- lat: double (nullable = true)
|-- long: double (nullable = true)
|-- name: string (nullable = true)
|-- station_id: integer (nullable = true)
```

```
status=status.drop('_id')
status.printSchema()
```

```
root
|-- bikes_available: integer (nullable = true)
|-- docks_available: integer (nullable = true)
|-- station_id: integer (nullable = true)
|-- time: string (nullable = true)
```

```
weather=weather.drop('_id')
weather.printSchema()
```

```
root
|-- cloud_cover: double (nullable = true)
|-- date: string (nullable = true)
|-- events: string (nullable = true)
|-- max_dew_point_f: string (nullable = true)
|-- max_gust_speed_mph: string (nullable = true)
|-- max_humidity: string (nullable = true)
|-- max_sea_level_pressure_inches: double (nullable = true)
|-- max_temperature_f: double (nullable = true)
|-- max_visibility_miles: string (nullable = true)
|-- max_wind_speed_mph: double (nullable = true)
|-- mean_dew_point_f: string (nullable = true)
|-- mean_humidity: string (nullable = true)
|-- mean_sea_level_pressure_inches: double (nullable = true)
|-- mean_temperature_f: double (nullable = true)
|-- mean_visibility_miles: string (nullable = true)
|-- mean_wind_speed_mph: double (nullable = true)
|-- min_dew_point_f: string (nullable = true)
|-- min_humidity: string (nullable = true)
|-- min_sea_level_pressure_inches: double (nullable = true)
|-- min_temperature_f: double (nullable = true)
|-- min_visibility_miles: string (nullable = true)
|-- precipitation_inches: string (nullable = true)
|-- wind_dir_degrees: double (nullable = true)
|-- zip_code: integer (nullable = true)
```

```
trip=trip.drop('_id')
trip.printSchema()
```

```
root
|-- bike_id: integer (nullable = true)
|-- duration: integer (nullable = true)
|-- end_date: string (nullable = true)
|-- end_station_id: integer (nullable = true)
|-- end_station_name: string (nullable = true)
|-- start_date: string (nullable = true)
|-- start_station_id: integer (nullable = true)
|-- start_station_name: string (nullable = true)
|-- subscription_type: string (nullable = true)
|-- trip_id: integer (nullable = true)
|-- zip_code: string (nullable = true)
```


3. Spark - Load Data from MongoDB

API Name	Memory	vCPUs	Instance Storage	Network Performance
m1.small	1.7 GiB	1 vCPUs	160 GiB HDD + 900MB swap	Low
m1.medium	3.75 GiB	1 vCPUs	410 GiB HDD	Moderate
m1.large	7.5 GiB	2 vCPUs	840 GiB (2 * 420 GiB HDD)	Moderate

What we learnt from the pipeline

Set up mongo-spark-connect

Make sure the version of the mongo-spark-connector is consistent with the spark version on the cluster!!

Failing this will result in the error when run the spark code.

Tip: to find the version of spark
`$ spark-submit --version`



```
Using Python version 2.7.12 (default, Nov 2 2017 19:20:38)
SparkSession available as 'spark'.
[>>> df = spark.read.format("com.mongodb.spark.sql.DefaultSource").option("uri", "mongodb://ec2-54-203-16-199.us-west-2.compute.amazonaws.com/bikeshare.weather").load()
18/01/17 21:43:35 WARN TaskSetManager: Lost task 0.0 in stage 0.0 (TID 0, 172.31.32.35): java.lang.NoSuchMethodError: org.apache.spark.sql.catalyst.analysis.TypeCoercion$.findTightestCommonType()Lscala.Function2;
    at com.mongodb.spark.sql.MongoInferSchema$.com$mongodb$spark$sql$MongoInferSchema$$compatibleType(MongoInferSchema.scala:135)
    at com.mongodb.spark.sql.MongoInferSchema$$anonfun$3$.apply(MongoInferSchema.scala:78)
    at com.mongodb.spark.sql.MongoInferSchema$$anonfun$3$.apply(MongoInferSchema.scala:78)
    at scala.collection.TraversableOnce$$anonfun$foldLeft$1$.apply(TraversableOnce.scala:157)
```

Spark version	connector	note
2.2.x	<code>org.mongodb.spark:mongo-spark-connector_2.11:2.2.0</code>	Usually on the Mac development machine
2.0.x	<code>org.mongodb.spark:mongo-spark-connector_2.11:2.0.0</code>	One the spark cluster set up by spark-ec2

FYR

Feature Selection

Station

Features:

Station ID, Station Name, latitude, longitude, number of Bikes available, city, install date

Trips

Features:

id, duration, start date, start station name, end date, end station name, bike id, subscription type, zip code, Number of Trip

Weather

Features:

temperature, dew point, humidity, sea level pressure, visibility, wind speed, gust speed, precipitation, cloud cover, weather condition, wind direction, zip code, date

- Join
- Generate New Features

Final Dataframe

Features:

Holiday , **Day of Week**, Station ID, Station Name, latitude, longitude, number of Bikes available, duration, start date, start station name, end date, end station name, bike id, subscription type, zip code, temperature, dew point, humidity, sea level pressure, visibility, wind speed, gust speed, precipitation, cloud cover, weather condition, wind direction, Number of Trips (label)

Label

```
tripRaw['start_date'] = tripRaw['start_date'].dt.floor('d')
```

```
tripAgg = tripRaw.groupby(['start_date', 'zip_code', 'start_station_name'],  
                           as_index=False)['id'].count()
```

```
tripAgg=tripAgg.rename(columns={'id': "num_trips"})
```

```
tripAgg.head()
```

	start_date	zip_code	start_station_name	num_trips
0	2013-08-29	10003	California Ave Caltrain Station	2
1	2013-08-29	10003	University and Emerson	2
2	2013-08-29	10009	Commercial at Montgomery	1
3	2013-08-29	10009	Mechanics Plaza (Market at Battery)	1
4	2013-08-29	10009	South Van Ness at Market	1

5. Processing

```
trip.show(5)
```

	_id bike_id duration	end_date end_station_id	end_station_name	id	start_date start
	_station_id	start_station_name	subscription_type	zip_code	
[5a5d4642b2c86795...]	520	63 8/29/2013 14:14	66 South Van Ness at...	4576 8/29/2013 14:13	
66 South Van Ness at...	Subscriber	94127			
[5a5d4642b2c86795...]	661	70 8/29/2013 14:43	10 San Jose City Hall	4607 8/29/2013 14:42	
10 San Jose City Hall	Subscriber	95138			
[5a5d4642b2c86795...]	48	71 8/29/2013 10:17	27 Mountain View Cit...	4130 8/29/2013 10:16	
27 Mountain View Cit...	Subscriber	97214			
[5a5d4642b2c86795...]	26	77 8/29/2013 11:30	10 San Jose City Hall	4251 8/29/2013 11:29	
10 San Jose City Hall	Subscriber	95060			
[5a5d4642b2c86795...]	527	103 8/29/2013 18:56	59 Golden Gate at Polk	4927 8/29/2013 18:54	
59 Golden Gate at Polk	Subscriber	94109			

only showing top 5 rows

```
df = targetDf.drop('_id', 'unnamed', 'date', 'installation_date', 'city')
```

```
df = df.na.drop()
```

```
#converting strings to numeric values
from pyspark.ml.feature import StringIndexer
```

```
def indexStringColumns(df, cols):
    #variable newdf will be updated several times
    newdf = df

    for c in cols:
        #For each given column, fits StringIndexerModel.
        si = StringIndexer(inputCol=c, outputCol=c+"-num")
        sm = si.fit(newdf)
        #Creates a DataFrame by putting the transformed values in the new column with suffix "-num"
        #and then drops the original columns.
        #and drop the "-num" suffix.
        newdf = sm.transform(newdf).drop(c)
        newdf = newdf.withColumnRenamed(c+"-num", c)

    return newdf
```

```
dfnumeric = indexStringColumns(df, ['start_station_name', 'events', 'dock_count'])
```

```
# Merging the data with Vector Assembler.
from pyspark.ml.feature import VectorAssembler
input_cols = ['zip_code', 'max_temperature_f', 'mean_temperature_f', 'min_temperature_f',
              'max_dew_point_f', 'mean_dew_point_f', 'min_dew_point_f', 'max_humidity', 'mean_humidity',
              'min_humidity', 'max_sea_level_pressure_inches', 'mean_sea_level_pressure_inches',
              'min_sea_level_pressure_inches', 'max_visibility_miles', 'mean_visibility_miles',
              'min_visibility_miles', 'max_wind_speed_mph', 'mean_wind_speed_mph',
              'max_gust_speed_mph', 'precipitation_inches', 'cloud_cover', 'wind_dir_degrees',
              'station_id', 'lat', 'long', 'holidays', 'day_of_week', 'start_station_name',
              'events', 'dock_count']
```

```
#VectorAssembler takes a number of column names(inputCols) and output column name (outputCol)
#and transforms a DataFrame to assemble the values in inputCols into one single vector with outputCol.
va = VectorAssembler(outputCol="features", inputCols=input_cols)
#1points - labeled data.
df_final = va.transform(dfnumeric).select("features", "label")
```

```
dfSets = df_final.randomSplit([0.8, 0.2], 1)
dfTrain = dfSets[0].cache()
dfTest = dfSets[1].cache()
```

Schema

Incorrect schema prevents model from fitting

Null Value

If data contains null value, model won't fit


5. Modeling & Results

```
dfSets = df_final.randomSplit([0.8, 0.2], 1)
dfTrain = dfSets[0].cache()
dfTest = dfSets[1].cache()
```

```
from pyspark.ml.regression import RandomForestRegressor
rf = RandomForestRegressor(maxDepth=20, maxBins=70)
rfmodel = rf.fit(dfTrain)
```

```
from pyspark.ml.evaluation import RegressionEvaluator
rfpredicts = rfmodel.transform(dfTest)
evaluator = RegressionEvaluator(labelCol="label", predictionCol="prediction", metricName="rmse")
rmse = evaluator.evaluate(rfpredicts)
print("RMSE = %d" % rmse)
```

RMSE = 1



```
rfpredicts.show(7)
```

features	label	prediction
[94041.0,56.0,49....]	1	1.6
[94041.0,58.0,48....]	1	1.55
[94041.0,60.0,51....]	1	1.9
[94041.0,61.0,50....]	1	1.4
[94041.0,61.0,56....]	2	1.8
[94041.0,63.0,51....]	1	1.8
[94041.0,63.0,56....]	2	1.8

only showing top 7 rows

6. Performance Evaluation - modeling.py pipeline



Read in data and
create spark dataframe

Numericalize string to
numbers

Assemble data into
vector and split into
training and test set

Fit the model and
report the result


- `time spark-submit modeling.py > output.txt`
final dataframe schema, baseline rmse
- real: wall clock time
- user: CPU time spent in user-mode code (outside the kernel) within the process
- Sys: CPU time spent in the kernel within the process
- Sys + user: how much actual CPU time the process used

6. Performance (local)

After time spark-submit modeling.py > output.txt
and finger crossed...

```
real    1m16.128s
user    1m42.300s
sys     0m2.709s
```

6. Spark-EC2 Cluster - 2 slaves (m1.large)

 **Spark Master at spark://ip-172-31-22-51.us-west-2.compute.internal:7077**

URL: spark://ip-172-31-22-51.us-west-2.compute.internal:7077
REST URL: spark://ip-172-31-22-51.us-west-2.compute.internal:6066 (cluster mode)
Alive Workers: 2
Cores in use: 4 Total, 0 Used
Memory in use: 12.1 GB Total, 0.0 B Used
Applications: 0 Running, 2 Completed
Drivers: 0 Running, 0 Completed
Status: ALIVE

Workers


Worker Id	Address	State	Cores	Memory
worker-20180118212347-172.31.26.153-52060	172.31.26.153:52060	ALIVE	2 (0 Used)	6.1 GB (0.0 B Used)
worker-20180118212347-172.31.26.73-47901	172.31.26.73:47901	ALIVE	2 (0 Used)	6.1 GB (0.0 B Used)

Running Applications

Application ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
----------------	------	-------	-----------------	----------------	------	-------	----------

Completed Applications

Application ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
app-20180118220930-0001	sparkml	4	5.8 GB	2018/01/18 22:09:30	root	FINISHED	2.0 min

 **Spark Worker at 172.31.26.153:52060**

ID: worker-20180118212347-172.31.26.153-52060
Master URL: spark://ip-172-31-22-51.us-west-2.compute.internal:7077
Cores: 2 (0 Used)
Memory: 6.1 GB (0.0 B Used)
[Back to Master](#)

Running Executors (0)

ExecutorID	Cores	State	Memory	Job Details	Logs
------------	-------	-------	--------	-------------	------

Finished Executors (3)

ExecutorID	Cores	State	Memory	Job Details	Logs
0	2	KILLED	5.8 GB	ID: app-20180118215852-0000 Name: sparkml User: root	stdout stderr

Spark-EC2

```
real    2m10.093s
user    1m5.436s
sys     0m5.008s
```

6. Performance Comparisons

Spark-EC2

```
real    2m10.093s
user    1m5.436s
sys     0m5.008s
```

Local

```
real    1m16.128s
user    1m42.300s
sys     0m2.709s
```

Time saved on AWS: **32 s**

6. Lesson learnt

1. Data size might change
2. Machine selection decision at different stage should be flexible
3. Never stuck in the idea of “big data” and ignore cost efficiency

The End

